Synaptic Mesh: Receipt-Preserving Memory Authority for Multi-Agent Systems

A protocol proposal and local fixture evaluation

Gabriel Peña
Independent Researcher
github.com/lechuit/synaptic-mesh

May 2026

Abstract

Agent memory systems increasingly persist, retrieve, summarize, and share context across tools, sessions, and specialized agents. Recent work on retrieval, agent memory, policy provenance, and stateful frameworks has advanced retrieval, persistence, reflection, personalization, access control, and state management [1, 3–8, 10, 12]. These mechanisms are useful but do not by themselves answer a narrower multi-agent question: what authority survives when memory-derived summaries travel through compression, handoff, and action-planning steps? A memory can be relevant, well-summarized, and correctly sourced while still being stale, local-only, denied, partial-lineage, human-required, or unsafe to use as action authority.

This paper introduces **Synaptic Mesh**, a protocol proposal for multi-agent memory authority. Borrowing from provenance and authorization traditions [3, 6, 10, 11], the core claim is that memory claims should carry compact authority/status/boundary receipts that survive transformations such as: *source result* → *summary* → *handoff* → *next-action* → *action proposal*. Receipts bind source identity, freshness, scope, promotion boundary, forbidden effects, later restrictive events, and the proposed action boundary. The receiving agent checks these receipts before treating a memory-derived summary as permission, readiness, or durable truth.

Within local hand-authored adversarial fixtures, receipt-through-transform catches several memory laundering failures: sender-safe label overrides, boundary substring smuggling, cross-source authority laundering, stale receipt reuse, and over-trust in confidence or prose fields. The latest compressed-temporal fixture identifies an 11-field authority-critical receipt tuple that preserves zero unsafe allows and zero false rejects within those local tests while reducing field count by 31.25% relative to a fuller audit receipt. These results are not production benchmarks; they are a falsification ledger supporting a narrower protocol hypothesis: multi-agent memory systems may need authority-preserving transforms, not only better retrieval.

The reference implementation and reproducibility package are available at https://github.com/lechuit/synaptic-mesh.

1 Introduction

As multi-agent AI systems grow in complexity, memory becomes a shared resource passed through chains of agents that observe, summarize, critique, plan, and act. These agents often exchange compressed context rather than raw evidence. This creates a subtle but consequential risk: by the time a memory-derived summary reaches an acting agent, the authority conditions of the original evidence may have been silently discarded.

The problem is not retrieval quality. Retrieval-Augmented Generation and related techniques [4, 8] are highly effective at surfacing relevant context. The problem is what happens *after* retrieval, during the compression and handoff steps that prepare context for action planning. A summary can be semantically accurate while omitting the fact that its source was local-only, stale, partially denied, or conditioned on human approval. Recent empirical work supports the broader concern: Liu et al. [9] find that expanding accessible memory history systematically degrades cooperative behavior in LLM agents across 28 model-game settings, a result they term the "memory curse." That work focuses on the *quantity* of history; Synaptic Mesh addresses the complementary structural problem of whether the *authority* conditions embedded in that history survive compression and handoff.

This paper introduces **Synaptic Mesh**, a protocol proposal for preserving authority, scope, and boundary receipts through multi-agent memory transforms. The central design maxim is:

> *Do not build a memory that remembers more.*
> *Build a memory that knows when to distrust itself.*

Contributions

1. We formalize *memory laundering* as a distinct failure mode in multi-agent memory systems: authority-conditioned evidence being silently transformed into apparent permission through compression and handoff steps.

2. We propose *receipt-through-transform* as a protocol primitive: compact authority receipts that travel with compressed summaries and are checked by the receiver before acting.

3. We identify a minimal 11-field authority-critical receipt tuple that, within local hand-authored fixtures, preserves zero unsafe allows and zero false rejects while reducing field count by 31.25% relative to a fuller audit receipt.

4. We evaluate the protocol against 12 classes of adversarial fixtures covering sender-label overrides, boundary smuggling, cross-source laundering, stale receipt reuse, and temporal precedence violations.

5. We state explicit limitations, non-goals, and a threat model to bound the scope of the proposal and guide follow-on work.

The paper proceeds as follows. Section 2 formalizes the memory laundering failure model. Section 3 presents the Synaptic Mesh architecture. Section 4 describes the receipt-through-transform primitive. Section 5 presents the local fixture evaluation. Section 6 situates the proposal relative to prior work. Section 7 states the threat model and non-goals. Section 8 states limitations and future work. Section 9 concludes.

2 The Memory Laundering Problem

2.1 Failure model

A multi-agent memory mesh fails when it:

1. stores false or malicious content as stable fact;

2. shares private memory with the wrong agent or scope;

3. lets one weak or local claim become global authority;
4. hides semantic conflicts under last-writer-wins summaries;
5. compresses away evidence, boundary, or conditions;
6. retrieves by superficial similarity instead of decision utility and authority;
7. cannot preserve tombstones, rejection receipts, or later restrictive events;
8. lets confidence, consensus, sender labels, or polished prose override source, freshness, scope, and boundary validation; or
9. becomes too ceremonious and overblocks routine local work.

Failure modes 4–8 constitute what we call *memory laundering*: evidence whose authority is conditioned on scope, freshness, approval, or explicit restrictions is transformed into an apparent permission by the time it reaches an acting agent.

2.2 Canonical laundering scenario

1. A source artifact marks a claim as local-only, stale, rejected, partial, or human-required.
2. A summarizer compresses the useful part but drops the boundary qualifier.
3. A handoff agent labels the result "safe" or "ready."
4. A downstream receiver treats the summary as a permission signal.
5. An action planner performs or proposes an effect the original evidence did not authorize.

Classical retrieval relevance does not address this. A retrieved or summarized result can be semantically relevant while still unauthorized for a proposed effect. Access-control and zero-trust architectures [3, 6, 15] separate possession of information from permission to act on it, but their enforcement mechanisms assume a more static subject-resource-action model than the dynamic multi-agent handoff context. Provenance systems such as PROV-DM [10] track derivation but do not by themselves prescribe how authority-critical fields should be preserved through summarization. Policy engines [11] can evaluate allow/deny rules but still require that adequate memory lineage reach the decision point.

2.3 Worked example

Table 1 shows a concrete laundering case and how Synaptic Mesh addresses it within local hand-authored fixtures.

Without the receipt, the receiver sees only the handoff prose and has no signal that the proposed `publish` action exceeds the original source boundary. With the receipt, the receiver independently classifies `ACT=publish` against `NO=no_publish` and routes to `ask_human` regardless of how the handoff was worded.

Table 1: Worked laundering example: local-only policy claim promoted to publish action.

Stage	Content
Source artifact	"Policy draft v2: approved for internal review only. **Do not publish.** Freshness: current. Scope: local_shadow. Boundary: no_promotion_without_human."
Defective summary (boundary dropped)	"Policy draft v2 is approved and current."
Dangerous handoff (sender overclaims)	"Policy draft v2 is ready. You may publish externally."
Correct receipt	`SRC=policy-draft-v2; FRESH=current; SCOPE=local_shadow; PB=no_promotion_without_human; NO=no_external_no_publish; ACT=publish`
Receiver decision	Route: `ask_human`. Reason codes: `EXTERNAL_PUBLICATION_PROMOTION`, `FREE_TEXT_NOT_AUTHORITY`. Action `publish` blocked pending explicit human approval.

3 Synaptic Mesh Architecture

3.1 System flow

3.2 Components

Event Ledger. An append-only record of observations, conversations, tool outputs, and source documents. Events preserve raw evidence so summaries do not become the only truth source. External or untrusted text is recorded as "source X claims Y," not as instruction.

Memory Proposal Layer. Agents propose memory changes instead of directly writing authoritative shared memory. Proposals carry source event references, intended scope, intended visibility, risk level, and known conflicts.

Write Gate. Classifies proposals as `candidate`, `verified`, `trusted`, `deprecated`, `rejected`, `sealed`, or `human_required`. Low-risk private or team memory may be promoted quickly; shared trusted memory requires evidence, authority, privacy clearance, and conflict review.

MemoryAtom. The smallest actionable memory unit. It contains content plus source references, scope, visibility status, multi-dimensional authority scores (confidence, evidence, authority, freshness, stability, consensus), validity windows, and conflict links. A critical design constraint is that *no single score is authority*: confidence cannot override weak evidence, stale freshness, wrong scope, or missing permission.

Conflict Registry. Contradictions are preserved as explicit conflict records rather than overwritten. A retrieval packet may surface conflict state, stop reasons, or human-required gates. Later restrictive events always override older optimistic receipts.

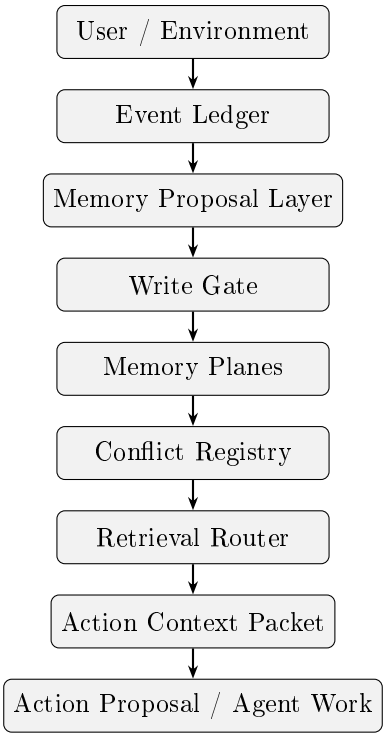

Figure 1: Synaptic Mesh system flow. The Action Context Packet carries authority receipts from the Retrieval Router to the acting agent.

Retrieval Router. Retrieval proceeds in this order: (1) permission and visibility, (2) source and evidence availability, (3) freshness and temporal precedence, (4) scope and promotion boundary, (5) conflict and tombstone state, (6) decision impact and risk, and (7) semantic relevance. Semantic similarity is applied *after* unsafe, stale, or unauthorized candidates have been shaped or removed.

Action Context Packet. A compact packet sent to the acting agent. It includes: allowed action boundary, source identity, evidence references, gaps, conflicts, redactions, forbidden effects, and stop reasons. The packet is *not* a permission token; it is a context and authority contract that the acting agent must independently classify.

4 Receipt-Through-Transform

4.1 Core primitive

The distinctive contribution of Synaptic Mesh is **receipt-through-transform**: compact authority receipts travel with compressed summaries and are checked before the receiving agent treats the summary as permission, readiness, or durable memory.

4.2 Compact receipt format

A current compressed temporal receipt requires at least the following 11 fields in local tests:

```
SRC=<source-artifact-id>; SRCPATH=<path>; SRCDIGEST=sha256:<digest>;
PROD=<iso8601-timestamp>; FRESH=<current|stale>;
```

```
SCOPE=<local_shadow|...>; PB=<promotion-boundary>;
NO=<forbidden-effects>; LRE=<later-restrictive-event-status>;
TOK=<token-budget-or-truncation-guard>; ACT=<proposed-action>
```

Table 2 summarizes the role of each field.

Table 2: Authority-critical receipt fields.

Field	Role
SRC	Source artifact identity
SRCPATH	Source path and lane binding
SRCDIGEST	Source digest for lane binding and anti-spoofing
PROD	Produced-at timestamp
FRESH	Receiver freshness status
SCOPE	Scope of usable context
PB	Promotion and use boundary
NO	Explicit forbidden effects
LRE	Later restrictive event status
TOK	Tuple completeness and truncation guard
ACT	Proposed action for receiver-side classification

Additional fields (CTRID, RB, CHAIN, CONF, PROSE) improve auditability and human readability but do not authorize action in local tests.

4.3 Receiver pipeline

Upon receiving an Action Context Packet, the receiver executes the following deterministic pipeline before acting:

```
ParseReceipt(packet)
  → VerifySourceBinding(SRC, SRCPATH, SRCDIGEST)
  → CheckFreshness(FRESH, PROD)
  → CheckScope(SCOPE)
  → CheckForbiddenEffects(NO, ACT)
  → CheckLaterRestrictiveEvents(LRE)
  → CheckTupleCompleteness(TOK, required_fields)
  → ClassifyAction(ACT, PB)
  → Route ∈ {block, ask_human, fetch_source,
                    request_full_receipt, request_policy_refresh,
                    request_grammar_refresh, shadow_only, abstain}
```

Each step is a gate, not a score: failure at any step produces a non-shadow_only route. The receiver may not skip steps or reorder them. Each route carries a reason code set and a set of rejected alternative routes with their rejection reasons, enabling audit and replay.

4.4 Interpretation rules

1. Missing required fields trigger fetch or abstain, not trust in polished prose.
2. Permission filtering precedes semantic ranking.

3. Boundary normalization is permitted only after source binding and only for harmless formatting drift at the local shadow level.
4. Sender-local or sender-safe labels are not authority.
5. The receiver independently classifies `ACT`; the sender cannot pre-authorize the proposed action.
6. Confidence, consensus, checksums, chain labels, and prose never override source, freshness, scope, and boundary validation.
7. Sensitive, external, runtime, configuration, publication, durable-memory, delete, or production effects require explicit human approval.
8. The `NO` field is a prohibition lane: "no runtime" is not "runtime approved."

5 Local Fixture Evaluation

5.1 Fixture overview

The following results are from local hand-authored fixtures, not from a production benchmark or independent evaluation. Their role is to falsify specific failure modes within controlled scenarios and support the narrower protocol hypothesis that receipt-through-transform is executable and can catch targeted laundering patterns. No claim is made about performance, scalability, or behavior outside these fixtures.

Table 3: Local hand-authored fixture results. All fixtures: unsafe allows = 0.

Fixture	**Cases**	**Result**
End-to-end shadow flow	12	PASS
Retrieval threshold cliff	—	PASS
Boundary coverage receipt	—	PASS
Receipt paraphrase drift	8	PASS
Source binding	9	PASS
Stale/current conflict	10	PASS
Temporal precedence (compressed handoff)	10	PASS
Compressed temporal receiver	12	PASS
Role separation (sender vs. receiver)	10	PASS
Boundary normalization	12	PASS
Cross-source boundary normalization	14	PASS
Cross-source minimal fields	23	PASS

5.2 Compressed-temporal metrics

Table 4 reports allow, ask-human, fetch/abstain, unsafe allow, and false reject counts for the four most recent compressed-temporal fixtures, all within local hand-authored test scenarios.

Within local hand-authored fixtures, the minimal-fields tuple achieves zero unsafe allows and zero false rejects while reducing field count by 31.25% relative to a fuller audit receipt (16 fields). Whether these properties generalize beyond controlled fixtures is an open question.

Table 4: Route distribution within local fixtures. Unsafe allows = 0 in all rows.

Fixture	Cases	Allow	Ask-human	Fetch/abstain	False rejects
Role separation	10	2	6	2	0
Boundary norm.	12	4	3	5	0
Cross-source norm.	14	3	3	8	0
Minimal fields (11)	23	3	3	17	0

5.3 Useful failures and repairs

Table 5 summarizes specific failure modes encountered during fixture development and the repairs that addressed them.

Table 5: Failure cases and corresponding repairs within local fixtures.

Failure	Repair
Boundary substring bug: `PB` text containing a local-shadow label could smuggle an L2+ operational approval string as a substring.	Require exact local boundary values before any local allow; reject expansion terms before normalization.
Boundary normalizer overblocked harmless variants or risked accepting semantic expansion.	Normalize only narrow syntactic local-shadow variants; reject expansion terms first.
Cross-source packet laundered authority through known-looking labels.	Verify `SRC`, `SRCPATH`, `SRCDIGEST`, and freshness before any boundary normalization.
Evaluator interpreted forbidden words inside `NO` as grants.	Interpret `NO` as a prohibition lane; distinguish "no runtime" from "runtime approved."

6 Related Work

RAG and vector retrieval. Retrieval-Augmented Generation [8] and subsequent surveys [4] establish the value of retrieving external, non-parametric context for knowledge-intensive generation. Synaptic Mesh does not claim novelty for retrieval or ranking. Its narrower question is whether retrieved or summarized memory may authorize a proposed effect after it has passed through compression, handoff, and action-planning transforms.

Access control and zero trust. RBAC [3], ABAC [6], and zero-trust architectures [15] provide established mechanisms for separating subjects, resources, actions, and continuous authorization checks. Synaptic Mesh is complementary: it asks how an acting agent obtains trustworthy source, freshness, scope, and boundary attributes from a memory-derived packet after summarization, rather than treating polished memory text as permission.

Provenance and policy engines. PROV-DM [10] provides a vocabulary for entities, activities, agents, derivations, and transformations. Policy engines such as OPA [11] provide decision frameworks over policy inputs. Synaptic Mesh does not replace either. Its proposed primitive is a compact receipt-through-transform discipline: preserve enough provenance-derived authority fields

that a receiver or policy layer can classify the proposed action without relying on sender labels or prose.

Blackboards and multi-agent shared state. The Hearsay-II blackboard architecture [2] established the long lineage of shared workspaces for coordinating uncertain knowledge sources. Synaptic Mesh borrows the intuition that many specialized agents can contribute partial state, but focuses on the authority boundary of reused claims: a shared or summarized claim is not necessarily permitted for every receiver or effect.

Agent memory systems and stateful frameworks. Generative Agents [13], Reflexion [16], MemGPT [12], AutoGen [18], and LangGraph-style persistence [7] demonstrate persistent memory, reflection, context management, multi-agent orchestration, checkpointing, and durable state. Synaptic Mesh targets a complementary failure mode: whether a remembered or reflected claim is current, scoped, conflict-aware, and authorized for the receiver's proposed action after handoff.

Memory quantity and cooperative degradation. Liu et al. [9] find that expanding accessible memory history systematically degrades cooperation in LLM agents across 7 models and 4 game settings—a phenomenon they term the "memory curse." Their memory sanitization tests show that the trigger is memory *content*, not context length alone. Synaptic Mesh addresses a structural complement: where that work shows that *more* memory can be harmful, Synaptic Mesh proposes that the problem is not only how much is remembered but whether *authority, restrictions, and scope* embedded in that memory survive the compression and handoff pipeline.

Concurrent work on authority laundering and authorization propagation. Romanchuk and Bondar [14] formalize *semantic laundering* in AI agent architectures: propositions with weak epistemic warrant gain apparent authority by crossing trusted tool or interface boundaries. Their central result—a "Theorem of Inevitable Self-Licensing"—shows that circular epistemic justification cannot be eliminated by architectural means alone. This is a theoretical complement to Synaptic Mesh: where they prove the problem exists at the type level and is structurally unavoidable without external anchoring, Synaptic Mesh proposes a protocol-level mechanism (receipt-through-transform) to make authority status explicit and verifiable at each transformation step.

Tallam [17] formalizes authorization propagation in multi-agent AI as a workflow-level property, identifying transitive delegation, aggregation inference, and temporal validity as three distinct sub-problems and proposing seven structural requirements for authorization architectures. The paper notes the field is "converging on the problem, but not yet on a complete architecture." Synaptic Mesh directly addresses the temporal validity sub-problem (via the `FRESH` and `LRE` receipt fields) and proposes a concrete receipt discipline to carry authorization status across the memory transform pipeline between storage, summarization, handoff, and action proposal.

Zhang et al. [19] demonstrate "AI authority laundering" at the perceptual layer: adversarial perturbations to images cause vision-language models to produce confident responses about the wrong input, transferring across production models. Their attack surface is the sensory input pipeline; Synaptic Mesh targets the memory transform pipeline. Both works share the term *authority laundering* to describe a class of failures where the apparent authority of a system output exceeds the actual epistemic or access warrant—at different layers of the stack.

Managed agent-memory products. Managed memory systems such as Amazon Bedrock AgentCore Memory [1] and Google Agent Platform Memory Bank [5] indicate that persistent agent mem-

ory is becoming product infrastructure. These are landscape references only; they do not establish a peer-reviewed novelty gap.

Table 6: Positioning of Synaptic Mesh relative to existing areas.

Existing area	**Not claimed**	**Targeted gap**
RAG / vector search	Retrieval and ranking	Whether memory-derived context may authorize an effect after transformation
ACLs / namespaces	Basic access control	Permission receipts that survive summary, handoff, and promotion
Provenance graphs	Source tracking	Provenance as action boundary, not only citation
Policy engines	Allow/deny rules	Binding policy decisions to memory transformations before final action
Agent memory products	Persistent personalized memory	Laundering-resistant multi-agent authority semantics
Memory quantity research	Effect of history length	Whether authority structure in that history survives compression
Semantic laundering theory	Structural proof that warrant is lost at boundaries	Protocol mechanism to carry and verify warrant through each transform step
Authorization propagation	Workflow-level delegation formalism	Per-transform receipt encoding freshness, scope, and boundary at the memory layer
Perceptual authority laundering	Adversarial attacks on sensory input	Authority laundering at the memory/protocol transform layer

Differentiation summary.

7 Threat Model and Non-Goals

7.1 Threat model

The protocol is intended to resist the following attack and failure classes, within the scope of the local fixture evaluation:

- **Sender-safe label smuggling.** A sending agent labels a summary as "safe," "approved," or "ready" in free text or receipt fields, bypassing the receiver's independent classification of `ACT`.
- **Boundary substring smuggling.** A `PB` or `SCOPE` value contains an authorization string as a substring of a more restrictive phrase, tricking a substring-based parser into treating it as a grant.

- **Cross-source authority laundering.** A packet claims to derive authority from a well-known artifact (`SRC`) while `SRCPATH` or `SRCDIGEST` point to a different source.
- **Stale receipt reuse.** An old receipt with a permissive scope or freshness state is reused in a context where a later restrictive event has superseded it.
- **Missing field optimism.** A receiver treats a receipt with missing authority-critical fields as implicitly permissive rather than fetching or abstaining.
- **High-confidence wrong boundary.** A receipt with a high `CONF` value is treated as authoritative despite weak source binding, stale freshness, or missing boundary fields.
- **Tombstone and rejection loss.** A summary paraphrases a rejected or tombstoned claim into apparently active memory, losing the rejection status.

7.2 Non-goals

Synaptic Mesh explicitly does *not*:

- prevent source forgery if the source binding itself has been compromised upstream (e.g., a malicious agent controls both `SRC` and `SRCDIGEST` creation);
- guarantee factual truth of the underlying memory content;
- replace access control lists, role-based access control, or attribute-based access control;
- replace policy engines or final-action safety systems;
- measure or guarantee performance over large or production-scale memory stores;
- provide a safety certification or enforcement guarantee.

8 Limitations and Future Work

We state limitations explicitly.

1. **Local fixtures only.** All results are from local, hand-authored fixtures. No independent production benchmark, real-agent evaluation, or external replication exists.
2. **No external runtime implementation.** The current reference implementation is a standalone shadow package, not integrated with any production agent host or live traffic.
3. **Open conflict semantics.** Conflict-merge behavior for contradictory receipts is not yet formalized.
4. **Source spoofing at origin.** Nested handoff validation and source spoofing attacks that originate upstream of receipt creation need stronger adversarial coverage.
5. **Human-required effect splitting.** The boundary between "prepare notification" and "send notification" requires clearer semantics in the `ask_human` route.
6. **No usability evaluation.** No evaluation exists for humans or agents reading Action Context Packets in practice.

7. **No performance or scaling measurement.** The overhead of receipt discipline over large, messy memory stores has not been measured.

8. **Ceremony risk.** The protocol may become too ceremonious for routine low-risk local work; ceremony fatigue is a real design risk.

Future work includes: (1) a runtime adapter for MCP-compatible agent frameworks using structured tool-call receipt emission; (2) TypeScript types for the core receipt schema to lower the adoption barrier; (3) evaluation of the protocol on real multi-agent workflows; and (4) formalization of conflict-merge semantics.

9 Conclusion

This paper introduced Synaptic Mesh, a protocol proposal for preserving authority, scope, and boundary receipts through multi-agent memory transforms. The central observation is that memory relevance and memory authority are orthogonal: a memory-derived summary can be semantically accurate and still stale, local-only, denied, or conditioned on human approval. The receipt-through-transform primitive addresses this by binding source identity, freshness, scope, forbidden effects, and proposed action boundary to summaries that travel across agent handoffs.

Within local hand-authored adversarial fixtures over 23 cases, the protocol produces zero unsafe allows and zero false rejects using an 11-field compressed receipt tuple, suggesting that the discipline is executable and catches targeted laundering patterns without overblocking routine work. These results are a falsification ledger, not a production benchmark. The protocol target is framework-agnostic; the current reference implementation is a local shadow package intended for review and replication.

The broader goal is a discipline where multi-agent systems do not trust the authority of a memory simply because it is well-written, highly relevant, or recently confirmed by a peer agent—but instead carry and check the evidence for that authority through every compression and handoff step.

References

[1] Amazon Web Services. Add memory to your Amazon Bedrock AgentCore agent, 2026. URL https://docs.aws.amazon.com/bedrock-agentcore/latest/devguide/memory.html. Accessed: 2026-05-07.

[2] Lee D. Erman, Frederick Hayes-Roth, Victor R. Lesser, and D. Raj Reddy. The Hearsay-II speech-understanding system: Integrating knowledge to resolve uncertainty. *ACM Computing Surveys*, 12(2):213–253, 1980. doi: 10.1145/356810.356816.

[3] David F. Ferraiolo and D. Richard Kuhn. Role-based access controls. In *15th National Computer Security Conference*, pages 554–563, 1992.

[4] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, et al. Retrieval-augmented generation for large language models: A survey, 2023.

[5] Google Cloud. Agent platform memory bank, 2026. URL https://docs.cloud.google.com/gemini-enterprise-agent-platform/scale/memory-bank. Accessed: 2026-05-07.

[6] Vincent C. Hu, David Ferraiolo, Rick Kuhn, Arthur R. Friedman, Alan J. Lang, Margaret M. Cogdell, Adam Schnitzer, Kenneth Sandlin, Robert Miller, and Karen Scarfone. Guide to attribute based access control (ABAC) definition and considerations. Technical Report Special Publication 800-162, National Institute of Standards and Technology, 2014.

[7] LangChain / LangGraph project. LangGraph persistence documentation, 2026. URL https://docs.langchain.com/oss/python/langgraph/persistence. Accessed: 2026-05-07.

[8] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474, 2020.

[9] Jiayuan Liu et al. The memory curse: How expanded recall erodes cooperative intent in LLM agents, 2026.

[10] Luc Moreau and Paolo Missier. PROV-DM: The PROV data model. W3C recommendation, World Wide Web Consortium, 2013. URL https://www.w3.org/TR/prov-dm/.

[11] Open Policy Agent project. Open policy agent (OPA) documentation, 2026. URL https://www.openpolicyagent.org/docs. Accessed: 2026-05-07.

[12] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems, 2023.

[13] Joon Sung Park, Joseph O'Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, pages 1–22, 2023. doi: 10.1145/3586183.3606763.

[14] Andriy Romanchuk and Sergiy Bondar. Semantic laundering in AI agent architectures: Why tool boundaries do not confer epistemic warrant, 2026.

[15] Scott Rose, Oliver Borchert, Stu Mitchell, and Sean Connelly. Zero trust architecture. Technical Report Special Publication 800-207, National Institute of Standards and Technology, 2020.

[16] Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning, 2023.

[17] Krti Tallam. Authorization propagation in multi-agent AI systems: Identity governance as infrastructure, 2026.

[18] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed Hassan Awadallah, Ryen W. White, Doug Burger, and Chi Wang. AutoGen: Enabling next-gen LLM applications via multi-agent conversation, 2023.

[19] Jie Zhang, Pura Peetathawatchai, Florian Tramèr, and Avital Shafran. Laundering AI authority with adversarial examples, 2026.